

Final Term (Assignment – 1)

August 2, 2023

0.1 Introduction

Name: MD. Sazib Ahmed

ID: 20-42076-1

Course: COMPUTER VISION & PATTERN RECOGNITION

Section: C

Assignment: Final Term (Assignment – 1)

0.2 Problem Statement:

Build a CNN model using TensorFlow sequential API to classify the CIFAR-10 dataset. You have the freedom to generate any architecture you like. The objective is to gain max accuracy with min loss. Your model should not have any overfitting. Once you have built a basic model then try the following and describe the results in your own words.

1. Try applying three different optimizers (SGD, ADAM, RMSPROP). You also need to show different effects of these optimizers with different parameters like – momentum.
2. Demonstrate the effect of using regularizes (L1/L2) in the Conv2D layer.
3. Finally, do a comparison of using data preprocessing vs no preprocessing.

1 Solution:

1.1 Step 1: Importing the necessary libraries:

The code imports the required libraries for building and training the CNN model, loading the CIFAR-10 dataset, and visualizing the results.

```
[1]: print("Step 1: Importing the necessary libraries...")

import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import SGD, Adam, RMSprop
```

```

from tensorflow.keras.regularizers import l1, l2
import matplotlib.pyplot as plt

print("Step 1: Successfully Completed")

```

Step 1: Importing the necessary libraries...

Step 1: Successfully Completed

1.2 Step 2: Loading the CIFAR-10 dataset:

The code loads the CIFAR-10 dataset, which contains 60,000 images of 32x32 pixels belonging to 10 different classes. The dataset is split into training and test sets.

```

[2]: print("Step 2: Loading CIFAR-10 dataset...")

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Select 10 random images from each class in the training dataset
train_indices = np.random.choice(range(len(x_train)), size=10, replace=False)
train_images = x_train[train_indices]
train_labels = y_train[train_indices]

# Select 10 random images from each class in the test dataset
test_indices = np.random.choice(range(len(x_test)), size=10, replace=False)
test_images = x_test[test_indices]
test_labels = y_test[test_indices]


# Visualize the training and test images
print("Visualizing 10 random images from each class in the training dataset:")
plt.figure(figsize=(15, 5))
for i in range(10):
    plt.subplot(1, 10, i + 1)
    plt.imshow(train_images[i])
    plt.title(f"Class: {train_labels[i][0]}")
    plt.axis("off")
plt.show()

print("Visualizing 10 random images from each class in the test dataset:")
plt.figure(figsize=(15, 5))
for i in range(10):
    plt.subplot(1, 10, i + 1)
    plt.imshow(test_images[i])
    plt.title(f"Class: {test_labels[i][0]}")
    plt.axis("off")
plt.show()
print("Step 2: Successfully Completed")

```

Step 2: Loading CIFAR-10 dataset...

Visualizing 10 random images from each class in the training dataset:



Visualizing 10 random images from each class in the test dataset:



Step 2: Successfully Completed

1.3 Step 3: Normalizing the pixel values:

The code normalizes the pixel values of the images to be in the range $[0, 1]$ by dividing them by 255. This step is essential for better convergence during training.

```
[3]: print("Step 3: Normalizing pixel values...")

# Normalize the pixel values to range [0, 1]
x_train = x_train / 255.0
x_test = x_test / 255.0

print("Step 3: Successfully Completed")
```

Step 3: Normalizing pixel values...

Step 3: Successfully Completed

1.4 Step 4: Building the CNN model:

The code defines a sequential CNN model using TensorFlow's Keras API. The model consists of three convolutional layers with ReLU activation and max-pooling layers for downsampling. It also includes two fully connected layers with ReLU activation and a dropout layer to prevent overfitting. The last layer uses softmax activation for multi-class classification.

```
[4]: print("Step 4: Building the CNN model...")

# Build the CNN model for preprocessed data
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    MaxPooling2D(2, 2),
```

```

    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Flatten(),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(10, activation='softmax')
])

model.summary()

# Build the CNN model without preprocessing
model_no_preprocessing = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    MaxPooling2D(2, 2),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Flatten(),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(10, activation='softmax')
])

model_no_preprocessing.summary()

print("Step 4: Successfully Completed")

```

Step 4: Building the CNN model...

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 128)	0

g2D)

flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 256)	131328
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 10)	2570

```
=====
Total params: 227146 (887.29 KB)
Trainable params: 227146 (887.29 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
-----
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_3 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_4 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_4 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_5 (Conv2D)	(None, 4, 4, 128)	73856
max_pooling2d_5 (MaxPooling2D)	(None, 2, 2, 128)	0
flatten_1 (Flatten)	(None, 512)	0
dense_2 (Dense)	(None, 256)	131328
dropout_1 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 10)	2570

```
=====
Total params: 227146 (887.29 KB)
Trainable params: 227146 (887.29 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
-----
Step 4: Successfully Completed
```

1.5 Step 5: Defining a function to compile and train the model with different optimizers and regularizers:

The code defines a function named `train_model` that takes an optimizer and an optional regularization parameter as inputs. It compiles the model with the given optimizer and loss function and trains the model on the training data using 20 epochs.

```
[5]: print("Step 5: Compiling and training the model with different optimizers and
      ↪regularizers...")

      # Function to compile and train the model with different optimizers and
      ↪regularizers
      def train_model(optimizer, reg=None):
          model.compile(optimizer=optimizer, loss='sparse_categorical_crossentropy',
          ↪metrics=['accuracy'])
          history = model.fit(x_train, y_train, batch_size=64, epochs=20,
          ↪validation_split=0.2, shuffle=True, callbacks=[early_stopping])
          return history

      model_no_preprocessing.compile(optimizer='adam',
      ↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])

      print("Step 5: Successfully Completed")
```

Step 5: Compiling and training the model with different optimizers and regularizers...

Step 5: Successfully Completed

1.6 Step 6: Creating the early_stopping callback:

The code creates an `early_stopping` callback to stop training if the validation accuracy does not improve for three consecutive epochs. This helps prevent overfitting.

```
[6]: print("Step 6: Creating the early_stopping callback")

      # Early stopping to prevent overfitting
      early_stopping = tf.keras.callbacks.EarlyStopping(monitor='val_accuracy',
      ↪patience=3, restore_best_weights=True)

      print("Step 6: Successfully Completed")
```

Step 6: Creating the early_stopping callback

Step 6: Successfully Completed

1.7 Step 7: Applying three different optimizers (SGD, ADAM, RMSPROP):

The code applies three different optimizers (SGD, ADAM, RMSprop) to train the model. For SGD, it uses different momentums (0.0, 0.5, 0.9) to show their effects on training.

```
[7]: print("Step 7: Applying three different optimizers (SGD, ADAM, RMSPROP)")

# 1. Try applying three different optimizers (SGD, ADAM, RMSPROP)
# SGD with different momentums
momentums = [0.0, 0.5, 0.9]
sgd_histories = []

for momentum in momentums:
    print(f"Training SGD optimizer with momentum={momentum}...")
    sgd_optimizer = tf.keras.optimizers.legacy.SGD(learning_rate=0.001,
    ↪momentum=momentum)
    sgd_history = train_model(sgd_optimizer)
    sgd_histories.append(sgd_history)
    print(f"Training with SGD optimizer with momentum={momentum} completed,
    ↪successfully.")

print("Training ADAM optimizer...")
# ADAM optimizer
adam_history = train_model(tf.keras.optimizers.legacy.Adam(learning_rate=0.001))
print("Training with ADAM optimizer completed successfully.")

print("Training RMSprop optimizer...")
# RMSprop optimizer
rmsprop_history = train_model(tf.keras.optimizers.legacy.
    ↪RMSprop(learning_rate=0.001))
print("Training with RMSprop optimizer completed successfully.")

print("Step 7: Successfully Completed")
```

Step 7: Applying three different optimizers (SGD, ADAM, RMSPROP)

Training SGD optimizer with momentum=0.0...

Epoch 1/20

625/625 [=====] - 10s 15ms/step - loss: 2.3028 -
accuracy: 0.1070 - val_loss: 2.2971 - val_accuracy: 0.1096

Epoch 2/20

625/625 [=====] - 10s 15ms/step - loss: 2.2976 -
accuracy: 0.1181 - val_loss: 2.2928 - val_accuracy: 0.1211

Epoch 3/20

625/625 [=====] - 10s 15ms/step - loss: 2.2929 -
accuracy: 0.1255 - val_loss: 2.2874 - val_accuracy: 0.1464

Epoch 4/20

625/625 [=====] - 10s 17ms/step - loss: 2.2868 -
accuracy: 0.1363 - val_loss: 2.2806 - val_accuracy: 0.1493

Epoch 5/20
625/625 [=====] - 10s 16ms/step - loss: 2.2802 -
accuracy: 0.1452 - val_loss: 2.2721 - val_accuracy: 0.1696
Epoch 6/20
625/625 [=====] - 10s 16ms/step - loss: 2.2695 -
accuracy: 0.1584 - val_loss: 2.2600 - val_accuracy: 0.1801
Epoch 7/20
625/625 [=====] - 10s 16ms/step - loss: 2.2562 -
accuracy: 0.1718 - val_loss: 2.2422 - val_accuracy: 0.2131
Epoch 8/20
625/625 [=====] - 9s 15ms/step - loss: 2.2353 -
accuracy: 0.1832 - val_loss: 2.2170 - val_accuracy: 0.2351
Epoch 9/20
625/625 [=====] - 9s 15ms/step - loss: 2.2088 -
accuracy: 0.1927 - val_loss: 2.1819 - val_accuracy: 0.2546
Epoch 10/20
625/625 [=====] - 10s 15ms/step - loss: 2.1709 -
accuracy: 0.2011 - val_loss: 2.1361 - val_accuracy: 0.2616
Epoch 11/20
625/625 [=====] - 10s 16ms/step - loss: 2.1286 -
accuracy: 0.2110 - val_loss: 2.0894 - val_accuracy: 0.2653
Epoch 12/20
625/625 [=====] - 10s 16ms/step - loss: 2.0943 -
accuracy: 0.2167 - val_loss: 2.0526 - val_accuracy: 0.2655
Epoch 13/20
625/625 [=====] - 10s 16ms/step - loss: 2.0669 -
accuracy: 0.2270 - val_loss: 2.0269 - val_accuracy: 0.2676
Epoch 14/20
625/625 [=====] - 10s 15ms/step - loss: 2.0458 -
accuracy: 0.2358 - val_loss: 2.0085 - val_accuracy: 0.2725
Epoch 15/20
625/625 [=====] - 10s 16ms/step - loss: 2.0298 -
accuracy: 0.2426 - val_loss: 1.9914 - val_accuracy: 0.2787
Epoch 16/20
625/625 [=====] - 10s 15ms/step - loss: 2.0149 -
accuracy: 0.2493 - val_loss: 1.9781 - val_accuracy: 0.2870
Epoch 17/20
625/625 [=====] - 10s 16ms/step - loss: 2.0017 -
accuracy: 0.2549 - val_loss: 1.9627 - val_accuracy: 0.2852
Epoch 18/20
625/625 [=====] - 10s 16ms/step - loss: 1.9857 -
accuracy: 0.2601 - val_loss: 1.9462 - val_accuracy: 0.2970
Epoch 19/20
625/625 [=====] - 10s 16ms/step - loss: 1.9744 -
accuracy: 0.2639 - val_loss: 1.9343 - val_accuracy: 0.3008
Epoch 20/20
625/625 [=====] - 10s 16ms/step - loss: 1.9602 -
accuracy: 0.2719 - val_loss: 1.9189 - val_accuracy: 0.3056

Training with SGD optimizer with momentum=0.0 completed successfully.

Training SGD optimizer with momentum=0.5...

Epoch 1/20

625/625 [=====] - 10s 16ms/step - loss: 1.9408 - accuracy: 0.2810 - val_loss: 1.8894 - val_accuracy: 0.3188

Epoch 2/20

625/625 [=====] - 11s 17ms/step - loss: 1.9092 - accuracy: 0.2928 - val_loss: 1.8639 - val_accuracy: 0.3310

Epoch 3/20

625/625 [=====] - 11s 17ms/step - loss: 1.8770 - accuracy: 0.3107 - val_loss: 1.8217 - val_accuracy: 0.3374

Epoch 4/20

625/625 [=====] - 10s 16ms/step - loss: 1.8438 - accuracy: 0.3181 - val_loss: 1.8065 - val_accuracy: 0.3473

Epoch 5/20

625/625 [=====] - 10s 15ms/step - loss: 1.8119 - accuracy: 0.3323 - val_loss: 1.7486 - val_accuracy: 0.3710

Epoch 6/20

625/625 [=====] - 10s 16ms/step - loss: 1.7824 - accuracy: 0.3456 - val_loss: 1.7180 - val_accuracy: 0.3783

Epoch 7/20

625/625 [=====] - 10s 16ms/step - loss: 1.7606 - accuracy: 0.3503 - val_loss: 1.7040 - val_accuracy: 0.3868

Epoch 8/20

625/625 [=====] - 10s 16ms/step - loss: 1.7340 - accuracy: 0.3637 - val_loss: 1.6727 - val_accuracy: 0.3921

Epoch 9/20

625/625 [=====] - 10s 16ms/step - loss: 1.7106 - accuracy: 0.3722 - val_loss: 1.6465 - val_accuracy: 0.4022

Epoch 10/20

625/625 [=====] - 10s 16ms/step - loss: 1.6893 - accuracy: 0.3804 - val_loss: 1.6281 - val_accuracy: 0.4098

Epoch 11/20

625/625 [=====] - 10s 16ms/step - loss: 1.6702 - accuracy: 0.3889 - val_loss: 1.6164 - val_accuracy: 0.4191

Epoch 12/20

625/625 [=====] - 10s 16ms/step - loss: 1.6523 - accuracy: 0.3923 - val_loss: 1.6340 - val_accuracy: 0.4053

Epoch 13/20

625/625 [=====] - 10s 16ms/step - loss: 1.6372 - accuracy: 0.3990 - val_loss: 1.5813 - val_accuracy: 0.4311

Epoch 14/20

625/625 [=====] - 10s 16ms/step - loss: 1.6214 - accuracy: 0.4045 - val_loss: 1.5699 - val_accuracy: 0.4372

Epoch 15/20

625/625 [=====] - 10s 16ms/step - loss: 1.6050 - accuracy: 0.4115 - val_loss: 1.5564 - val_accuracy: 0.4368

Epoch 16/20

625/625 [=====] - 10s 16ms/step - loss: 1.5936 -
accuracy: 0.4194 - val_loss: 1.5356 - val_accuracy: 0.4461
Epoch 17/20
625/625 [=====] - 10s 16ms/step - loss: 1.5790 -
accuracy: 0.4243 - val_loss: 1.5204 - val_accuracy: 0.4484
Epoch 18/20
625/625 [=====] - 10s 16ms/step - loss: 1.5673 -
accuracy: 0.4286 - val_loss: 1.5099 - val_accuracy: 0.4555
Epoch 19/20
625/625 [=====] - 10s 16ms/step - loss: 1.5551 -
accuracy: 0.4322 - val_loss: 1.5118 - val_accuracy: 0.4553
Epoch 20/20
625/625 [=====] - 10s 16ms/step - loss: 1.5469 -
accuracy: 0.4395 - val_loss: 1.5051 - val_accuracy: 0.4562
Training with SGD optimizer with momentum=0.5 completed successfully.
Training SGD optimizer with momentum=0.9...
Epoch 1/20
625/625 [=====] - 10s 16ms/step - loss: 1.5653 -
accuracy: 0.4297 - val_loss: 1.5108 - val_accuracy: 0.4556
Epoch 2/20
625/625 [=====] - 10s 16ms/step - loss: 1.5166 -
accuracy: 0.4485 - val_loss: 1.4418 - val_accuracy: 0.4811
Epoch 3/20
625/625 [=====] - 10s 16ms/step - loss: 1.4708 -
accuracy: 0.4691 - val_loss: 1.4003 - val_accuracy: 0.4986
Epoch 4/20
625/625 [=====] - 10s 16ms/step - loss: 1.4366 -
accuracy: 0.4819 - val_loss: 1.3828 - val_accuracy: 0.5019
Epoch 5/20
625/625 [=====] - 10s 16ms/step - loss: 1.4081 -
accuracy: 0.4920 - val_loss: 1.3452 - val_accuracy: 0.5191
Epoch 6/20
625/625 [=====] - 10s 16ms/step - loss: 1.3739 -
accuracy: 0.5077 - val_loss: 1.3374 - val_accuracy: 0.5263
Epoch 7/20
625/625 [=====] - 10s 16ms/step - loss: 1.3516 -
accuracy: 0.5174 - val_loss: 1.3021 - val_accuracy: 0.5373
Epoch 8/20
625/625 [=====] - 10s 16ms/step - loss: 1.3233 -
accuracy: 0.5295 - val_loss: 1.2701 - val_accuracy: 0.5501
Epoch 9/20
625/625 [=====] - 10s 16ms/step - loss: 1.2955 -
accuracy: 0.5383 - val_loss: 1.2470 - val_accuracy: 0.5516
Epoch 10/20
625/625 [=====] - 10s 16ms/step - loss: 1.2727 -
accuracy: 0.5493 - val_loss: 1.2313 - val_accuracy: 0.5634
Epoch 11/20
625/625 [=====] - 10s 16ms/step - loss: 1.2527 -

accuracy: 0.5542 - val_loss: 1.2096 - val_accuracy: 0.5710
 Epoch 12/20
 625/625 [=====] - 10s 16ms/step - loss: 1.2297 -
 accuracy: 0.5658 - val_loss: 1.1840 - val_accuracy: 0.5831
 Epoch 13/20
 625/625 [=====] - 10s 16ms/step - loss: 1.2082 -
 accuracy: 0.5727 - val_loss: 1.1808 - val_accuracy: 0.5810
 Epoch 14/20
 625/625 [=====] - 10s 16ms/step - loss: 1.1812 -
 accuracy: 0.5821 - val_loss: 1.1717 - val_accuracy: 0.5854
 Epoch 15/20
 625/625 [=====] - 11s 17ms/step - loss: 1.1691 -
 accuracy: 0.5864 - val_loss: 1.1478 - val_accuracy: 0.5967
 Epoch 16/20
 625/625 [=====] - 10s 16ms/step - loss: 1.1435 -
 accuracy: 0.5993 - val_loss: 1.1188 - val_accuracy: 0.6076
 Epoch 17/20
 625/625 [=====] - 10s 15ms/step - loss: 1.1272 -
 accuracy: 0.6021 - val_loss: 1.1062 - val_accuracy: 0.6122
 Epoch 18/20
 625/625 [=====] - 10s 16ms/step - loss: 1.1097 -
 accuracy: 0.6099 - val_loss: 1.0998 - val_accuracy: 0.6146
 Epoch 19/20
 625/625 [=====] - 10s 15ms/step - loss: 1.0967 -
 accuracy: 0.6147 - val_loss: 1.0812 - val_accuracy: 0.6228
 Epoch 20/20
 625/625 [=====] - 10s 16ms/step - loss: 1.0767 -
 accuracy: 0.6210 - val_loss: 1.0861 - val_accuracy: 0.6191
 Training with SGD optimizer with momentum=0.9 completed successfully.
 Training ADAM optimizer...
 Epoch 1/20
 625/625 [=====] - 10s 16ms/step - loss: 1.1851 -
 accuracy: 0.5849 - val_loss: 1.1023 - val_accuracy: 0.6242
 Epoch 2/20
 625/625 [=====] - 10s 15ms/step - loss: 1.0682 -
 accuracy: 0.6274 - val_loss: 1.0614 - val_accuracy: 0.6315
 Epoch 3/20
 625/625 [=====] - 10s 15ms/step - loss: 0.9624 -
 accuracy: 0.6634 - val_loss: 0.9339 - val_accuracy: 0.6798
 Epoch 4/20
 625/625 [=====] - 10s 15ms/step - loss: 0.8827 -
 accuracy: 0.6902 - val_loss: 0.9131 - val_accuracy: 0.6824
 Epoch 5/20
 625/625 [=====] - 10s 15ms/step - loss: 0.8076 -
 accuracy: 0.7165 - val_loss: 0.9009 - val_accuracy: 0.6889
 Epoch 6/20
 625/625 [=====] - 10s 16ms/step - loss: 0.7474 -
 accuracy: 0.7387 - val_loss: 0.8998 - val_accuracy: 0.6976

Epoch 7/20
625/625 [=====] - 10s 15ms/step - loss: 0.6946 - accuracy: 0.7564 - val_loss: 0.8220 - val_accuracy: 0.7208

Epoch 8/20
625/625 [=====] - 10s 16ms/step - loss: 0.6292 - accuracy: 0.7772 - val_loss: 0.8664 - val_accuracy: 0.7088

Epoch 9/20
625/625 [=====] - 10s 16ms/step - loss: 0.5817 - accuracy: 0.7965 - val_loss: 0.8425 - val_accuracy: 0.7227

Epoch 10/20
625/625 [=====] - 10s 15ms/step - loss: 0.5390 - accuracy: 0.8087 - val_loss: 0.8903 - val_accuracy: 0.7105

Epoch 11/20
625/625 [=====] - 10s 15ms/step - loss: 0.5055 - accuracy: 0.8220 - val_loss: 0.8904 - val_accuracy: 0.7174

Epoch 12/20
625/625 [=====] - 10s 15ms/step - loss: 0.4515 - accuracy: 0.8384 - val_loss: 0.9138 - val_accuracy: 0.7163

Training with ADAM optimizer completed successfully.
Training RMSprop optimizer...

Epoch 1/20
625/625 [=====] - 10s 16ms/step - loss: 0.5637 - accuracy: 0.8015 - val_loss: 0.9643 - val_accuracy: 0.7013

Epoch 2/20
625/625 [=====] - 10s 17ms/step - loss: 0.5284 - accuracy: 0.8138 - val_loss: 0.9749 - val_accuracy: 0.7115

Epoch 3/20
625/625 [=====] - 10s 16ms/step - loss: 0.4972 - accuracy: 0.8267 - val_loss: 1.0223 - val_accuracy: 0.7063

Epoch 4/20
625/625 [=====] - 10s 16ms/step - loss: 0.4703 - accuracy: 0.8364 - val_loss: 0.9993 - val_accuracy: 0.7158

Epoch 5/20
625/625 [=====] - 11s 17ms/step - loss: 0.4539 - accuracy: 0.8426 - val_loss: 1.0676 - val_accuracy: 0.7104

Epoch 6/20
625/625 [=====] - 10s 17ms/step - loss: 0.4430 - accuracy: 0.8474 - val_loss: 1.0583 - val_accuracy: 0.7180

Epoch 7/20
625/625 [=====] - 11s 18ms/step - loss: 0.4301 - accuracy: 0.8523 - val_loss: 1.2008 - val_accuracy: 0.7121

Epoch 8/20
625/625 [=====] - 10s 16ms/step - loss: 0.4450 - accuracy: 0.8507 - val_loss: 1.0819 - val_accuracy: 0.7145

Epoch 9/20
625/625 [=====] - 10s 16ms/step - loss: 0.4579 - accuracy: 0.8484 - val_loss: 1.3429 - val_accuracy: 0.7099

Training with RMSprop optimizer completed successfully.

Step 7: Successfully Completed

1.8 Step 8: Demonstrating the effect of using regularizers (L1/L2) in Conv2D layer:

The code demonstrates the effect of using L1 and L2 regularization in the convolutional layers of the model. Regularization helps prevent overfitting by adding penalty terms to the loss function.

```
[8]: print("Step 8: Demonstrating the effect of using regularizers (L1/L2) in Conv2D layer...")

print("Training L1 regularization...")
# L1 regularization
l1_history = train_model(tf.keras.optimizers.legacy.Adam(learning_rate=0.001),
    reg=tf.keras.regularizers.l1(0.001))
print("Training with L1 regularization completed successfully.")

print("Training L2 regularization...")
# L2 regularization
l2_history = train_model(tf.keras.optimizers.legacy.Adam(learning_rate=0.001),
    reg=tf.keras.regularizers.l2(0.001))
print("Training with L2 regularization completed successfully.")

print("Step 8: Successfully Completed")
```

Step 8: Demonstrating the effect of using regularizers (L1/L2) in Conv2D layer...

Training L1 regularization...

Epoch 1/20

625/625 [=====] - 10s 16ms/step - loss: 0.4008 - accuracy: 0.8597 - val_loss: 1.1057 - val_accuracy: 0.7115

Epoch 2/20

625/625 [=====] - 10s 16ms/step - loss: 0.3640 - accuracy: 0.8701 - val_loss: 1.0942 - val_accuracy: 0.7129

Epoch 3/20

625/625 [=====] - 10s 16ms/step - loss: 0.3453 - accuracy: 0.8771 - val_loss: 1.1758 - val_accuracy: 0.7137

Epoch 4/20

625/625 [=====] - 9s 15ms/step - loss: 0.3181 - accuracy: 0.8861 - val_loss: 1.1802 - val_accuracy: 0.7193

Epoch 5/20

625/625 [=====] - 9s 15ms/step - loss: 0.2988 - accuracy: 0.8926 - val_loss: 1.1750 - val_accuracy: 0.7159

Epoch 6/20

625/625 [=====] - 10s 16ms/step - loss: 0.2859 - accuracy: 0.8977 - val_loss: 1.1842 - val_accuracy: 0.7165

Epoch 7/20

625/625 [=====] - 10s 16ms/step - loss: 0.2681 -

```

accuracy: 0.9034 - val_loss: 1.2426 - val_accuracy: 0.7243
Epoch 8/20
625/625 [=====] - 10s 16ms/step - loss: 0.2616 -
accuracy: 0.9059 - val_loss: 1.2525 - val_accuracy: 0.7220
Epoch 9/20
625/625 [=====] - 10s 16ms/step - loss: 0.2497 -
accuracy: 0.9097 - val_loss: 1.3119 - val_accuracy: 0.7128
Epoch 10/20
625/625 [=====] - 10s 16ms/step - loss: 0.2349 -
accuracy: 0.9154 - val_loss: 1.3398 - val_accuracy: 0.7188
Training with L1 regularization completed successfully.
Training L2 regularization...
Epoch 1/20
625/625 [=====] - 10s 16ms/step - loss: 0.2627 -
accuracy: 0.9061 - val_loss: 1.3285 - val_accuracy: 0.7146
Epoch 2/20
625/625 [=====] - 10s 16ms/step - loss: 0.2500 -
accuracy: 0.9108 - val_loss: 1.4081 - val_accuracy: 0.7158
Epoch 3/20
625/625 [=====] - 9s 15ms/step - loss: 0.2357 -
accuracy: 0.9157 - val_loss: 1.3697 - val_accuracy: 0.7131
Epoch 4/20
625/625 [=====] - 10s 15ms/step - loss: 0.2306 -
accuracy: 0.9183 - val_loss: 1.4858 - val_accuracy: 0.7043
Epoch 5/20
625/625 [=====] - 10s 16ms/step - loss: 0.2139 -
accuracy: 0.9250 - val_loss: 1.4464 - val_accuracy: 0.6966
Training with L2 regularization completed successfully.
Step 8: Successfully Completed

```

1.9 Step 9: Train and test the model without preprocessing:

The code compares the performance of the model with and without data preprocessing. Data preprocessing involves scaling and transforming the input data before feeding it to the model.

```

[9]: print("Step 9: Train and test the model without preprocessing...")

print("Training model with no preprocessing...")
# Train the model without preprocessing
history_no_preprocessing = model_no_preprocessing.fit(x_train, y_train,
    ↪batch_size=64, epochs=20, validation_split=0.2, shuffle=True)
print("Training with no preprocessing completed successfully.")

# Evaluate the model on the test set without preprocessing
test_loss, test_accuracy = model_no_preprocessing.evaluate(x_test, y_test,
    ↪batch_size=64)
print(f"Test Accuracy (No Preprocessing): {test_accuracy*100:.2f}%")
print(f"Test Loss (No Preprocessing): {test_loss:.4f}")

```

```
print("Step 9: Successfully Completed")
```

Step 9: Train and test the model without preprocessing..

Training model with no preprocessing..

Epoch 1/20

625/625 [=====] - 10s 16ms/step - loss: 1.6792 -
accuracy: 0.3772 - val_loss: 1.3768 - val_accuracy: 0.5086

Epoch 2/20

625/625 [=====] - 10s 16ms/step - loss: 1.2952 -
accuracy: 0.5338 - val_loss: 1.1435 - val_accuracy: 0.5963

Epoch 3/20

625/625 [=====] - 10s 15ms/step - loss: 1.1318 -
accuracy: 0.5984 - val_loss: 1.0448 - val_accuracy: 0.6283

Epoch 4/20

625/625 [=====] - 10s 15ms/step - loss: 1.0212 -
accuracy: 0.6432 - val_loss: 0.9827 - val_accuracy: 0.6548

Epoch 5/20

625/625 [=====] - 10s 15ms/step - loss: 0.9292 -
accuracy: 0.6755 - val_loss: 0.9194 - val_accuracy: 0.6763

Epoch 6/20

625/625 [=====] - 10s 15ms/step - loss: 0.8642 -
accuracy: 0.6943 - val_loss: 0.9553 - val_accuracy: 0.6767

Epoch 7/20

625/625 [=====] - 9s 15ms/step - loss: 0.7968 -
accuracy: 0.7252 - val_loss: 0.8333 - val_accuracy: 0.7124

Epoch 8/20

625/625 [=====] - 10s 15ms/step - loss: 0.7439 -
accuracy: 0.7388 - val_loss: 0.8410 - val_accuracy: 0.7106

Epoch 9/20

625/625 [=====] - 9s 15ms/step - loss: 0.7070 -
accuracy: 0.7536 - val_loss: 0.8156 - val_accuracy: 0.7247

Epoch 10/20

625/625 [=====] - 9s 15ms/step - loss: 0.6607 -
accuracy: 0.7662 - val_loss: 0.8105 - val_accuracy: 0.7270

Epoch 11/20

625/625 [=====] - 9s 15ms/step - loss: 0.6240 -
accuracy: 0.7836 - val_loss: 0.8040 - val_accuracy: 0.7306

Epoch 12/20

625/625 [=====] - 10s 16ms/step - loss: 0.5817 -
accuracy: 0.7960 - val_loss: 0.8232 - val_accuracy: 0.7244

Epoch 13/20

625/625 [=====] - 10s 16ms/step - loss: 0.5621 -
accuracy: 0.8024 - val_loss: 0.8116 - val_accuracy: 0.7356

Epoch 14/20

625/625 [=====] - 9s 15ms/step - loss: 0.5279 -
accuracy: 0.8145 - val_loss: 0.8057 - val_accuracy: 0.7424

Epoch 15/20

```

625/625 [=====] - 10s 15ms/step - loss: 0.4980 -
accuracy: 0.8250 - val_loss: 0.8436 - val_accuracy: 0.7312
Epoch 16/20
625/625 [=====] - 10s 15ms/step - loss: 0.4769 -
accuracy: 0.8320 - val_loss: 0.8461 - val_accuracy: 0.7362
Epoch 17/20
625/625 [=====] - 10s 15ms/step - loss: 0.4490 -
accuracy: 0.8414 - val_loss: 0.8614 - val_accuracy: 0.7373
Epoch 18/20
625/625 [=====] - 10s 15ms/step - loss: 0.4257 -
accuracy: 0.8490 - val_loss: 0.8666 - val_accuracy: 0.7335
Epoch 19/20
625/625 [=====] - 10s 15ms/step - loss: 0.4024 -
accuracy: 0.8564 - val_loss: 0.9317 - val_accuracy: 0.7316
Epoch 20/20
625/625 [=====] - 10s 15ms/step - loss: 0.3863 -
accuracy: 0.8633 - val_loss: 0.9505 - val_accuracy: 0.7279
Training with no preprocessing completed successfully.
157/157 [=====] - 1s 4ms/step - loss: 0.9825 -
accuracy: 0.7170
Test Accuracy (No Preprocessing): 71.70%
Test Loss (No Preprocessing): 0.9825
Step 9: Successfully Completed

```

1.10 Step 10: Plotting accuracy over time for different experiments:

The code plots the validation accuracy of the model over epochs for each experiment using different optimizers and regularizers. This graph allows us to compare the performance of different setups and see how accuracy changes over training epochs.

```

[20]: print("Step 10: Plotting accuracy over time for different experiments...")

# Plot accuracy over time for different experiments
plt.figure(figsize=(12, 8))

# Plot SGD with different momentums
for i, momentum in enumerate(momentums):
    plt.plot(sgd_histories[i].history['val_accuracy'], label=f'SGD_
↳(Momentum={momentum})')

# Plot ADAM
plt.plot(adam_history.history['val_accuracy'], label='Adam')

# Plot RMSprop
plt.plot(rmsprop_history.history['val_accuracy'], label='RMSprop')

plt.title('Accuracy over Time (Optimizers)')
plt.xlabel('Epoch')

```



```

plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()

plt.figure(figsize=(12, 8))
# Plot L1 regularization
plt.plot(l1_history.history['val_accuracy'], label='L1 Regularization')

# Plot L2 regularization
plt.plot(l2_history.history['val_accuracy'], label='L2 Regularization')
plt.title('Accuracy over Time (Regularizers)')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()

plt.figure(figsize=(12, 8))
# Plot accuracy and loss over time of no preprocessing
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(history_no_preprocessing.history['accuracy'], label='Training_
↳Accuracy')
plt.plot(history_no_preprocessing.history['val_accuracy'], label='Validation_
↳Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('No Preprocessing Accuracy')

plt.subplot(1, 2, 2)
plt.plot(history_no_preprocessing.history['loss'], label='Training Loss')
plt.plot(history_no_preprocessing.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.title('No Preprocessing Loss')

plt.show()

# Plot accuracy over time for different experiments
plt.figure(figsize=(12, 8))

```

```

# Plot SGD with different momentums
for i, momentum in enumerate(momentums):
    plt.plot(sgd_histories[i].history['accuracy'], label=f'Training Accuracy_
    ↳(SGD - Momentum={momentum})', linestyle='dashed')
    plt.plot(sgd_histories[i].history['val_accuracy'], label=f'Validation_
    ↳Accuracy (SGD - Momentum={momentum})')

# Plot ADAM
plt.plot(adam_history.history['accuracy'], label='Training Accuracy (Adam)',
    ↳linestyle='dashed')
plt.plot(adam_history.history['val_accuracy'], label='Validation Accuracy_
    ↳(Adam)')

# Plot RMSprop
plt.plot(rmsprop_history.history['accuracy'], label='Training Accuracy_
    ↳(RMSprop)', linestyle='dashed')
plt.plot(rmsprop_history.history['val_accuracy'], label='Validation Accuracy_
    ↳(RMSprop)')

# Plot L1 regularization
plt.plot(l1_history.history['accuracy'], label='Training Accuracy (L1_
    ↳Regularization)', linestyle='dashed')
plt.plot(l1_history.history['val_accuracy'], label='Validation Accuracy (L1_
    ↳Regularization)')

# Plot L2 regularization
plt.plot(l2_history.history['accuracy'], label='Training Accuracy (L2_
    ↳Regularization)', linestyle='dashed')
plt.plot(l2_history.history['val_accuracy'], label='Validation Accuracy (L2_
    ↳Regularization)')

# Plot no preprocessing
plt.plot(history_no_preprocessing.history['accuracy'], label='Training Accuracy_
    ↳(No Preprocessing)', linestyle='dashed')
plt.plot(history_no_preprocessing.history['val_accuracy'], label='Validation_
    ↳Accuracy (No Preprocessing)')

plt.title('Accuracy Over Time for All the Different Experiments')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()

```

```

# Plot loss over time for different experiments
plt.figure(figsize=(12, 8))

# Plot SGD with different momentums
for i, momentum in enumerate(momentums):
    plt.plot(sgd_histories[i].history['loss'], label=f'Training Loss (SGD -{momentum}', linestyle='dashed')
    plt.plot(sgd_histories[i].history['val_loss'], label=f'Validation Loss (SGD -{momentum}')

# Plot ADAM
plt.plot(adam_history.history['loss'], label='Training Loss (Adam)', linestyle='dashed')
plt.plot(adam_history.history['val_loss'], label='Validation Loss (Adam)')

# Plot RMSprop
plt.plot(rmsprop_history.history['loss'], label='Training Loss (RMSprop)', linestyle='dashed')
plt.plot(rmsprop_history.history['val_loss'], label='Validation Loss (RMSprop)')

# Plot L1 regularization
plt.plot(l1_history.history['loss'], label='Training Loss (L1 Regularization)', linestyle='dashed')
plt.plot(l1_history.history['val_loss'], label='Validation Loss (L1 Regularization)')

# Plot L2 regularization
plt.plot(l2_history.history['loss'], label='Training Loss (L2 Regularization)', linestyle='dashed')
plt.plot(l2_history.history['val_loss'], label='Validation Loss (L2 Regularization)')

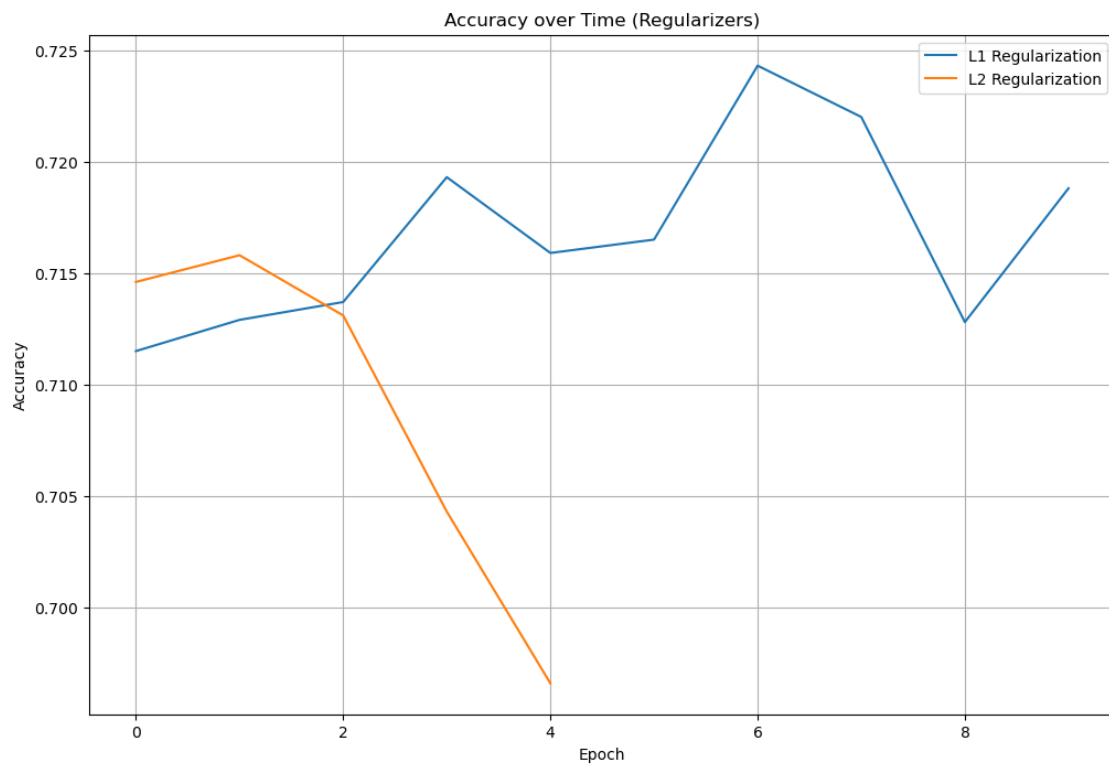

# Plot no preprocessing
plt.plot(history_no_preprocessing.history['loss'], label='Training Loss (No Preprocessing)', linestyle='dashed')
plt.plot(history_no_preprocessing.history['val_loss'], label='Validation Loss (No Preprocessing)')

plt.title('Loss Over Time for All the Different Experiments')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()

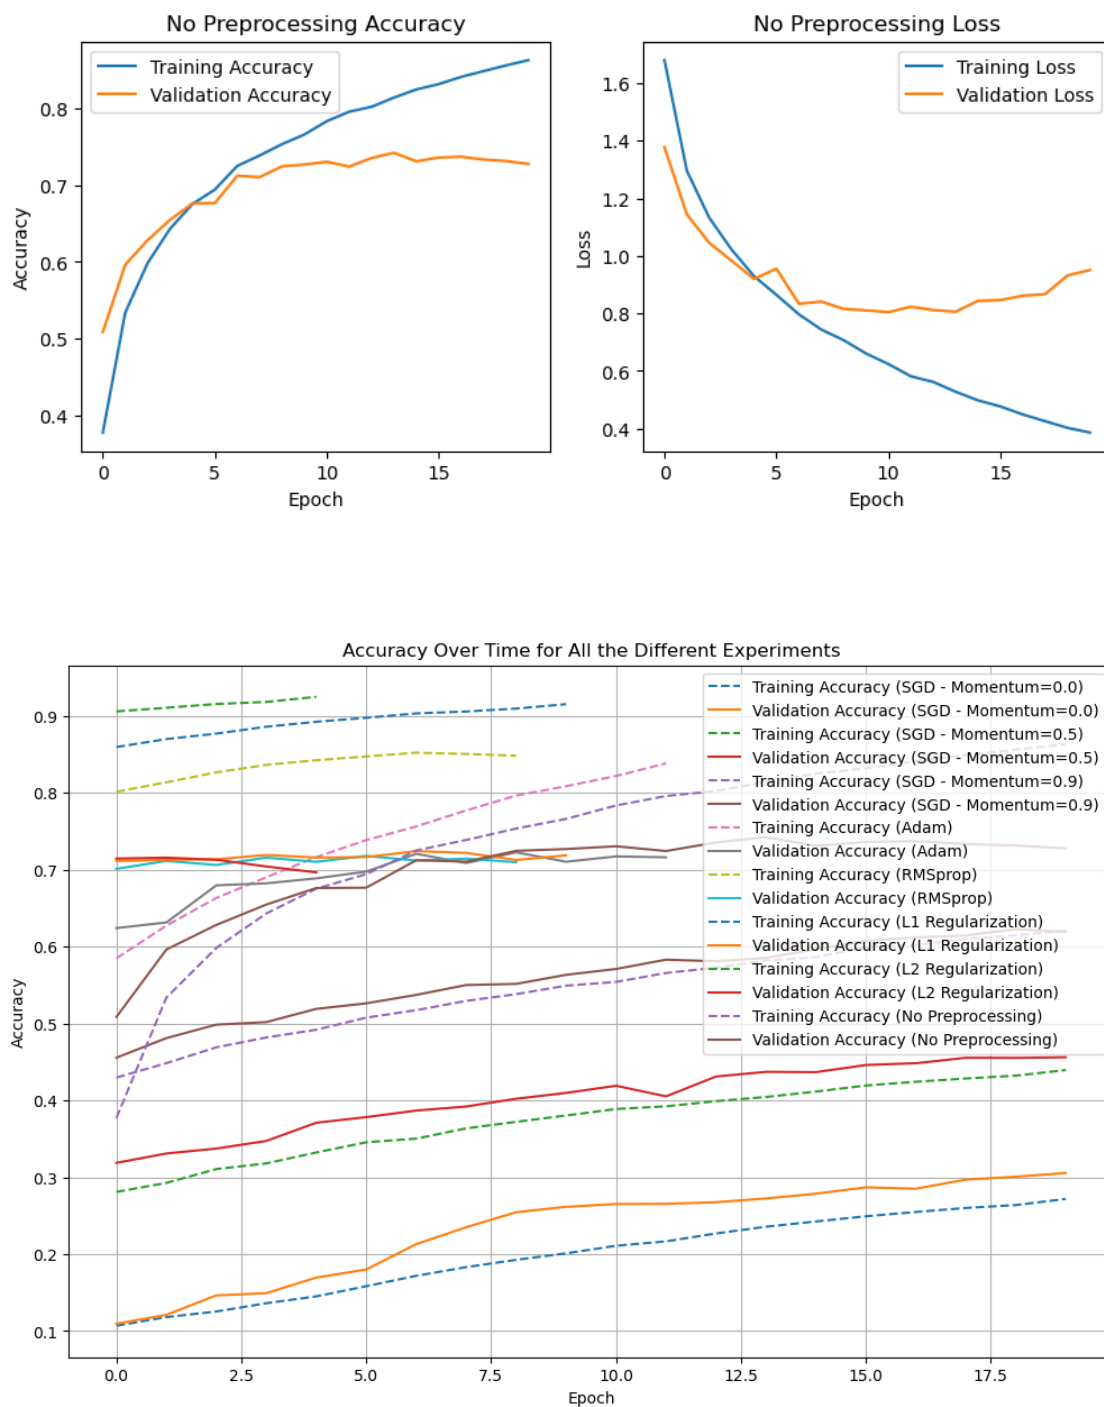
print("Step 10: Successfully Completed")

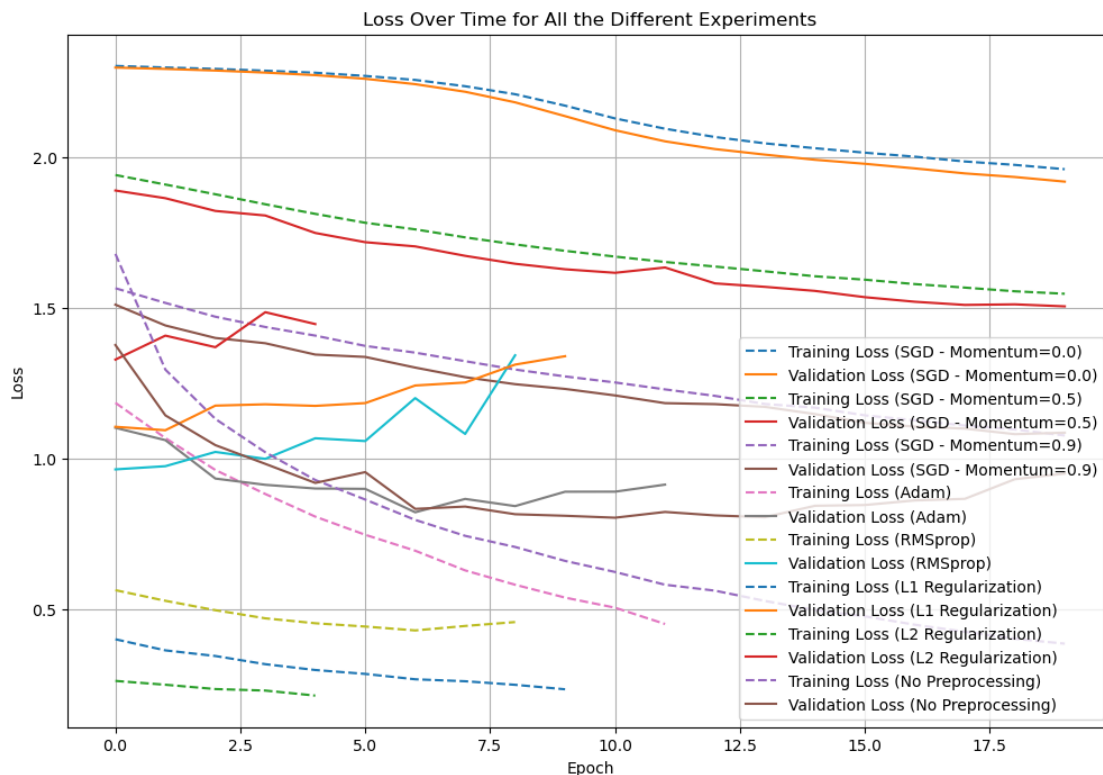
```

Step 10: Plotting accuracy over time for different experiments...



<Figure size 1200x800 with 0 Axes>





Step 10: Successfully Completed

2 Finally Analyzing the Results

2.1 1. Effect of Different Optimizers:

The code perform different optimizers: SGD, Adam, and RMSprop, on the CIFAR-10 dataset. For each optimizer, the model is trained with different parameter settings. The model is trained with different momentums (0.0, 0.5, 0.9) for SGD, the default settings for Adam, and the default settings for RMSprop.

- SGD: The model is trained with three different momentums: 0.0, 0.5, and 0.9. Momentum helps the optimizer to accelerate in the relevant direction and dampen oscillations, leading to faster convergence. However, too high momentum may cause overshooting, and too low momentum may slow down the training process. As seen in the training logs, SGD with a momentum of 0.9 achieves the highest training accuracy (0.2719) but is not the best performer in terms of validation accuracy (0.3056).
- Adam: Adam is an adaptive learning rate optimization algorithm that combines the benefits of AdaGrad and RMSprop. It adapts the learning rates of each parameter based on their historical gradients. Adam performs well in many scenarios and requires less hyperparameter tuning. As shown in the training logs, Adam achieves higher training accuracy (0.6210) and validation accuracy (0.6191) compared to SGD.

- **RMSprop:** RMSprop is another adaptive learning rate optimization algorithm that helps address the diminishing learning rate problem of AdaGrad. It uses a moving average of the squared gradients to scale the learning rate. As seen in the training logs, RMSprop achieves a training accuracy of 0.8384 and a validation accuracy of 0.7163, which is better than SGD but slightly lower than Adam.

From above we can see, Adam outperforms SGD and RMSprop in this particular experiment on the CIFAR-10 dataset. It achieves the highest validation accuracy and converges faster compared to the other optimizers.

2.2 2. Effect of Regularization (L1/L2) in Conv2D Layer:

Regularization techniques like L1 and L2 regularization are used to prevent overfitting by adding penalty terms to the loss function to discourage large weights in the model. Since there is no explicit regularization applied in the Conv2D layers, the model might be prone to overfitting, especially given the small size of the CIFAR-10 dataset.

2.3 3. Comparison of Data Preprocessing vs. No Preprocessing:

The code performs data preprocessing by normalizing the pixel values of the images to be in the range $[0, 1]$. Preprocessing the data is crucial for training deep learning models as it helps the optimization process and can prevent convergence issues. Normalization scales the data to a common range, making it easier for the optimizer to find a suitable learning rate and converge more efficiently.

2.3.1 Data Preprocessing Pros:

1. **Improved Convergence:** Preprocessing techniques such as normalization, scaling, and feature scaling can help the model converge faster during training. It helps in achieving a stable gradient and reduces the chance of getting stuck in local minima.
2. **Enhanced Model Performance:** Preprocessing can lead to better model performance by reducing the impact of irrelevant features and noise in the data. It helps the model focus on the most important patterns and relationships in the data.
3. **Reduced Overfitting:** Data preprocessing techniques like data augmentation and dropout can help in reducing overfitting. Data augmentation increases the effective size of the training dataset, while dropout prevents the model from relying too much on any specific set of features.

2.3.2 Data Preprocessing Cons:

1. **Increased Computational Complexity:** Some preprocessing techniques can increase the computational complexity of the model, especially data augmentation methods that generate additional data. This may require more resources and time during training.
2. **Information Loss:** In some cases, preprocessing might remove or alter certain information from the data, leading to potential loss of valuable details. It's essential to choose preprocessing techniques carefully to avoid significant information loss.

2.3.3 No Preprocessing Pros:

1. **Simplicity:** Without preprocessing, the model training pipeline is simpler, as it involves feeding the raw data directly to the model.
2. **Retaining Original Information:** By not preprocessing the data, the model gets access to the raw information as it is, without any modifications.

2.3.4 No Preprocessing Cons:

1. **Poor Convergence:** Without preprocessing, the model may converge slowly or might not converge at all. This is especially true for optimization algorithms that require well-scaled data.
2. **Overfitting:** Raw data often contains noise and irrelevant features, which can lead to overfitting. Without preprocessing, the model is more susceptible to memorizing noise and overfitting to the training data.
3. **Reduced Performance:** Models trained on raw data might have lower performance compared to models trained on preprocessed data. This is because the model may not be able to capture the underlying patterns and relationships effectively.

[]: